

Distributed Computing and Introduction to High Performance Computing

Imad Kissami¹

¹Mohammed VI Polytechnic University, Benguerir, Morocco



Outline of this lecture

Point-to-point Communications

- Blocking Send and Receive
- Simultaneous Send and Receive

Point-to-point Communications

Blocking Send MPI_SEND

```
1 MPI_SEND(buf, count, datatype, dest, tag, comm, code)
2
3 <type>:: buf
4 integer :: count, datatype
5 integer :: dest, tag, comm, code
```

- ➡ Sending, from the address `buf`, a message of `count` elements of type `datatype`, tagged `tag`, to the process of rank `dest` in the communicator `comm`.
- ➡ the execution remains **blocked** until the message can be re-written without risk of overwriting the value to be sent. In other words, the execution is blocked as long as the message has not been received.

```
1 COMM.Send(self, data, int dest, int tag=0)
```

Point-to-point Communications

Blocking Receive MPI_RECV

```
1 MPI_RECV(buf, count, datatype, source, tag, comm, status_msg, code)
2
3 <type>:: buf
4 integer :: count, datatype
5 integer :: source, tag, comm, code
6 integer, dimension(MPI_STATUS_SIZE) :: status_msg
```

- ➡ Receiving, at the address buf, a message of count elements of type datatype, tagged tag, from the process of rank source in the communicator comm.
- ➡ status_msg stores the state of a receive operation : source, tag, code, ...
- ➡ An MPI_RECV can only be associated to an MPI_SEND if these two calls have the same envelope (source, dest, tag, comm).
- ➡ the execution remains **blocked** until the message content corresponds to the received message.

```
1 data = COMM.recv(self, source, int tag=0)
2 # or
3 Comm.Recv(self, buf, int source, int tag=0, Status status=None)
```

Point-to-point Communications

Blocking Send / Receive Full example

```
1 from mpi4py import MPI
2
3 COMM = MPI.COMM_WORLD
4 RANK = COMM.Get_rank()
5
6 tag = 99
7
8
9 if RANK == 2:
10     sendbuf = 1000
11     COMM.send(sendbuf, dest=5, tag=tag)
12
13 if RANK == 5:
14     recvbuf = COMM.recv(source=2, tag=tag)
15     print("I, process 5, I received ",recvbuf," from the process 2.")
```

`mpirun -n 6 python sendrecv.py`

I, process 5, I received 1000 from the process 2.

Point-to-point Communications

Simultaneous send and receive MPI_SENDRECV

```
1 MPI_SENDRECV( sendbuf , sendcount , sendtype ,  
2               dest , sendtag ,  
3               recvbuf , recvcount , recvtype ,  
4               source , recvtag , comm , status_msg , code )  
5  
6 <type> :: sendbuf , recvbuf  
7 integer :: sendcount , recvcount  
8 integer :: sendtype , recvtype  
9 integer :: source , dest , sendtag , recvtag , comm , code  
10  
11 integer , dimension( MPI_STATUS_SIZE ) :: status_msg
```

- ➡ Sending, from the address sendbuf, a message of sendcount elements of type sendtype, tagged sendtag, to the process dest in the communicator comm ;
- ➡ Receiving, at the address recvbuf, a message of recvcount elements of type recvtype, tagged recvtag, from the process source in the communicator comm.
- ➡ Here, the receiving zone recvbuf must be different from the sending zone sendbuf.

```
1 Comm.Sendrecv( self , sendbuf , int dest , int sendtag=0 , recvbuf=None , int source=ANY_SOURCE , int recvtag=ANY_TAG , Status status=None )
```

Point-to-point Communications

Simultaneous send and receive MPI_SENDRECV: Full example

```
1  from mpi4py import MPI
2
3  COMM = MPI.COMM_WORLD
4  RANK = COMM.Get_rank()
5  SIZE = COMM.Get_size()
6  tag = 99
7
8  #We define the process we will communicate with (we suppose that we have exactly ↔
   2 processes)
9
10 num_proc = (RANK+1)%2
11
12 sendbuf = RANK + 1000
13 recvdata = COMM.sendrecv(sendbuf, num_proc, sendtag=tag, recvtag = tag, status=↔
   None)
14 # or
15 recvdata = COMM.sendrecv(sendbuf, num_proc)
16
17 print("I, process {proc_send}, I received {data} from the process "
18       "{proc_recv}.".format(proc_send = RANK, data = recvdata, proc_recv = ↔
   num_proc))
```

```
mpirun -n 2 python simultaneoussendrecv.py
```

I, process 0, I received 1001 from the process 1.

I, process 1, I received 1000 from the process 0.

Point-to-point Communications

Simultaneous send and receive MPI_SENDRECV: Remarks

In the case of a synchronous implementation of the MPI_SEND() subroutine, if we replace the MPI_SENDRECV() subroutine in the example above by MPI_SEND() followed by MPI_RECV(), the code will deadlock. Indeed, each of the two processes will wait for a receipt confirmation, which will never come because the two sending operations would stay suspended.

```
1 COMM.send(sendbuf, dest=5, tag=tag)
2 recvbuf = COMM.recv(source=2, tag=tag)
```