

Distributed Computing and Introduction to High Performance Computing

Imad Kissami¹

¹Mohammed VI Polytechnic University, Benguerir, Morocco



Derived datatypes

Introduction

- In communications, exchanged data have different datatypes : `MPI_INTEGER`, `MPI_REAL`, `MPI_COMPLEX`, etc.
- We can create more complex data structures by using subroutines such as
- `MPI_TYPE_CONTIGUOUS()`, `MPI_TYPE_VECTOR()`, `MPI_TYPE_INDEXED()` or `MPI_TYPE_CREATE_STRUCT()`
- Derived datatypes allow exchanging non-contiguous or non-homogenous data in the memory and limiting the number of calls to communications subroutines

Derived datatypes




Figure: Hierarchy of the MPI constructors

Derived datatypes

Contiguous datatypes : MPI_TYPE_CONTIGUOUS()

- MPI_TYPE_CONTIGUOUS() creates a data structure from a homogenous set of existing datatypes contiguous in memory.

```
1 MPI_TYPE_CONTIGUOUS(count, old_type, new_type, code)
2
3 integer, intent(in) :: count, old_type
4 integer, intent(out) :: new_type, code
```

```
1 Datatype.Create_contiguous(self, int count)
```

Derived datatypes

MPI_TYPE_COMMIT and MPI_TYPE_FREE()

- Before using a new derived datatype, it is necessary to validate it with the MPI_TYPE_COMMIT() subroutine.

```
1 MPI_TYPE_COMMIT(new_type, code)
2
3 integer, intent(inout) :: new_type
4 integer, intent(out) :: code
```

```
1 new_type.Commit()
```

- The freeing of a derived datatype is made by using the MPI_TYPE_FREE() subroutine.

```
1 MPI_TYPE_FREE(new_type, code)
2 integer, intent(inout) :: new_type
3 integer, intent(out) :: code
```

```
1 new_type.Free()
```

Derived datatypes

Contiguous datatypes : MPI_TYPE_CONTIGUOUS()

■ Example :

```
1  count = 5
2  size = 10
3
4  type_ligne = MPI.DOUBLE.Create_contiguous(count)
5  type_ligne.Commit()
6
7  if rank==0:
8      data = np.array([i for i in range(size)], dtype=np.float64)
9      comm.Send([data,1,type_ligne],dest=1)
10     print("Original data",data, rank)
11
12  elif rank == 1:
13      data = -1*np.ones(size, dtype=np.float64)
14      comm.Recv([data,1,type_ligne],source=0)
15      print("Received data",data, rank)
16
17  type_ligne.Free()
```

`mpirun -n 2 python3 create_contiguous.py`

Original data [0. 1. 2. 3. 4. 5. 6. 7. 8. 9.] 0

Received data [0. 1. 2. 3. 4. -1. -1. -1. -1. -1.] 1

Derived datatypes

Constant stride : MPI_TYPE_VECTOR()

- MPI_TYPE_VECTOR() creates a data structure from a homogenous set of existing datatypes separated by a constant stride in memory. The stride is given in number of elements.

```
1 MPI_TYPE_VECTOR(count, block_length, stride, old_type, new_type, code)
2
3 integer, intent(in) :: count, block_length
4 integer, intent(in) :: stride
5 integer, intent(in) :: old_type
6
7 integer, intent(out) :: new_type, code
```

```
1 Datatype.Create_vector(self, int count, int blocklength, int stride)
```

Derived datatypes

Constant stride : MPI_TYPE_VECTOR()

■ Example :

```
1  stride = 2
2  count = 5
3  blocklen = 1
4  size = 10
5
6  type_colum = MPI.DOUBLE.Create_vector(count, blocklen, stride)
7  type_colum.Commit()
8
9  if rank==0:
10     data = np.array([i for i in range(size)], dtype=np.float64)
11     comm.Send([data, 1, type_colum], dest=1)
12     print("Original data", data, rank)
13 elif rank==1:
14     data = -1*np.ones(size, dtype=np.float64)
15     comm.Recv([data, 1, type_colum], source=0)
16     print("Received data", data, rank)
17
18 type_colum.Free()
```

```
mpirun -n 2 python3 create_vector.py
```

```
Original data [0. 1. 2. 3. 4. 5. 6. 7. 8. 9.] 0
```

```
Received data [ 0. -1.  2. -1.  4. -1.  6. -1.  8. -1.] 1
```


Derived datatypes

Constant stride : MPI_TYPE_CREATE_HVECTOR()

- MPI_TYPE_CREATE_HVECTOR() creates a data structure from a homogenous set of existing datatype separated by a constant stride in memory. The stride is given in bytes.
- This call is useful when the old type is no longer a base datatype (MPI_INTEGER, MPI_REAL,...) but a more complex datatype constructed by using MPI subroutines, because in this case the stride can no longer be given in number of elements.

```
1 MPI_TYPE_CREATE_HVECTOR(count, block_length, stride, old_type, ↵  
    new_type, code)  
2  
3 integer, intent(in) :: count, block_length  
4 integer(kind=MPI_ADDRESS_KIND), intent(in) :: stride  
5 integer, intent(in) :: old_type  
6  
7 integer, intent(out) :: new_type, code
```

```
1 Datatype.Create_hvector(self, int count, int blocklength, Aint stride)
```

Derived datatypes

Homogenous datatypes of variable strides

- `MPI_TYPE_INDEXED()` allows creating a data structure composed of a sequence of blocks containing a variable number of elements separated by a variable stride in memory. The stride is given in number of elements.
- `MPI_TYPE_CREATE_HINDEXED()` has the same functionality as `MPI_TYPE_INDEXED()` except that the strides separating two data blocks are given in bytes. This subroutine is useful when the old datatype is not an MPI base datatype (`MPI_INTEGER`, `MPI_REAL`, ...). We cannot therefore give the stride in number of elements of the old datatype.
- For `MPI_TYPE_CREATE_HINDEXED()`, as for `MPI_TYPE_CREATE_HVECTOR()`, use `MPI_TYPE_SIZE()` or `MPI_TYPE_GET_EXTENT()` in order to obtain in a portable way the size of the stride in bytes.

Derived datatypes

Homogenous datatypes of variable strides : `MPI_TYPE_INDEXED()`

`nb=3, blocks_lengths=(2,1,3), displacements=(0,3,7)`



Figure: The `MPI_TYPE_INDEXED` constructor

```
1 MPI_TYPE_INDEXED(nb, block_lengths, displacements, old_type, new_type, ↵  
    code)  
2  
3 integer, intent(in) :: nb  
4 integer, intent(in), dimension(nb) :: block_lengths  
5 integer, intent(in), dimension(nb) :: displacements  
6 integer, intent(in) :: old_type  
7  
8 integer, intent(out) :: new_type, code
```

```
1 Datatype.Create_indexed(self, blocklengths, displacements)
```

Derived datatypes

Homogenous datatypes of variable strides : `MPI_TYPE_INDEXED()`

■ Example :

```
1  size = 10
2  count = 3
3
4  counts = [2, 1, 3]
5  displacements = [0, 3, 7]
6
7  indexedtype = MPI.INT64_T.Create_indexed(counts, displacements)
8  indexedtype.Commit()
9
10 if rank==0:
11     data = np.array([i for i in range(size)], dtype=np.float64)
12     comm.Send([data,1,indexedtype],dest=1)
13     print("Original data",data, rank)
14 elif rank==1:
15     data = -1*np.ones(size, dtype=np.float64)
16     comm.Recv([data,1,indexedtype],source=0)
17     print("Received data",data, rank)
18
19 indexedtype.Free()
```

```
mpirun -n 2 python3 create_indexed.py
```

```
Original data [0. 1. 2. 3. 4. 5. 6. 7. 8. 9.] 0
```

```
Received data [ 0.  1. -1.  3. -1. -1. -1.  7.  8.  9.] 1
```

Derived datatypes

Homogenous datatypes of variable strides : `MPI_TYPE_CREATE_HINDEXED()`




Figure: The `MPI_TYPE_CREATE_HINDEXED` constructor

```
1 MPI_TYPE_CREATE_HINDEXED(nb, block_lengths, displacements, old_type ←  
    , new_type, code)  
2  
3 integer, intent(in) :: nb  
4 integer, intent(in), dimension(nb) :: block_lengths  
5 integer(kind=MPI_ADDRESS_KIND), intent(in), dimension(nb) :: ←  
    displacements  
6 integer, intent(in) :: old_type  
7  
8 integer, intent(out) :: new_type, code
```

```
1 Datatype.Create_hindexed(self, blocklengths, displacements)
```

Derived datatypes

Homogenous datatypes of variable strides : `MPI_TYPE_INDEXED()`

■ Example : triangular matrix

In the following example, each of the two processes :

1. Initializes its matrix (positive growing numbers on process 0 and negative decreasing numbers on process 1).
2. Constructs its datatype : triangular matrix (superior for the process 0 and inferior for the process 1).
3. Sends its triangular matrix to the other process and receives back a triangular matrix which it stores in the same place which was occupied by the sent matrix.
4. Frees its resources.

Derived datatypes

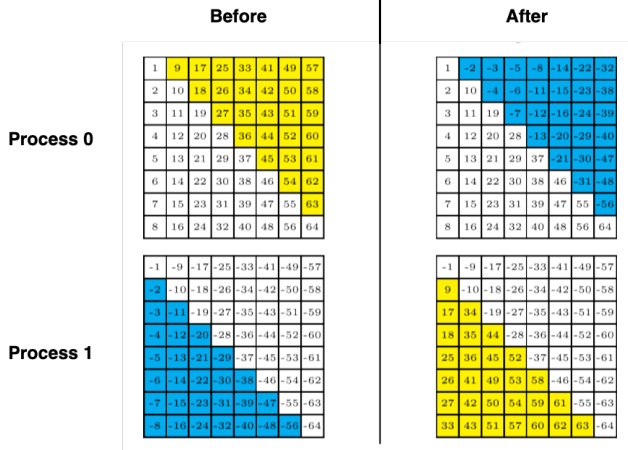


Figure: – Exchange between the two processes

Derived datatypes

Homogenous datatypes of variable strides : `MPI_TYPE_INDEXED()`

■ Example :

```
1 from mpi4py import MPI
2 import numpy as np
3 comm = MPI.COMM_WORLD
4 nb_procs = comm.Get_size()
5 rank = comm.Get_rank()
6 n = 8; sign = -1
7 if rank == 0: sign = 1
8
9 a = [sign*i for i in range(1,n*n+1,1)]
10 Matrix = np.array(a)
11 Matrix = np.reshape(Matrix, (n,n)).transpose()
12
13 if rank == 0:
14     displacements = [n*i for i in range(n)]
15     block_lengths = [i for i in range(n)]
16 else:
17     displacements = [n*i+i+1 for i in range(n)]
18     block_lengths = [n-i-1 for i in range(n)]
19
20 type_triangle = MPI.DOUBLE.Create_indexed(block_lengths, displacements)
21 type_triangle.Commit()
22
23 num_proc = (rank+1)%2
24 comm.Send([Matrix,1,type_triangle],dest=num_proc)
25 comm.Recv([Matrix,1,type_triangle],source=num_proc)
26
27 type_triangle.Free()
28 print(Matrix, rank)
```


Derived datatypes

Homogenous datatypes of variable strides : `MPI_TYPE_INDEXED()`

- Example : Matrix after permutation

```
mpirun -n 2 python3 matrixExchange.py
```

```
[[ 1 -2 -3 -5 -8 -14 -22 -32]
 [ 2 10 -4 -6 -11 -15 -23 -38]
 [ 3 11 19 -7 -12 -16 -24 -39]
 [ 4 12 20 28 -13 -20 -29 -40]
 [ 5 13 21 29 37 -21 -30 -47]
 [ 6 14 22 30 38 46 -31 -48]
 [ 7 15 23 31 39 47 55 -56]
 [ 8 16 24 32 40 48 56 64]] 0
```

```
[[ -1 -9 -17 -25 -33 -41 -49 -57]
 [ 9 -10 -18 -26 -34 -42 -50 -58]
 [ 17 34 -19 -27 -35 -43 -51 -59]
 [ 18 35 44 -28 -36 -44 -52 -60]
 [ 25 36 45 52 -37 -45 -53 -61]
 [ 26 41 49 53 58 -46 -54 -62]
 [ 27 42 50 54 59 61 -55 -63]
 [ 33 43 51 57 60 62 63 -64]] 1
```

Communicators

Introduction

The purpose of communicators is to create subgroups on which we can carry out operations such as collective or point-to-point communications. Each subgroup will have its own communication space.

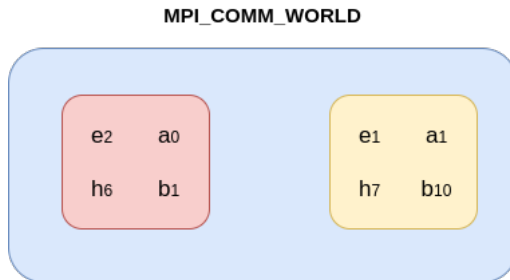


Figure: Communicator partitioning

Communicators

Example

For example, we want to broadcast a collective message to even-ranked processes and another message to odd-ranked processes.

- Looping on send/recv can be very detrimental especially if the number of processes is high. Also a test inside the loop would be compulsory in order to know if the sending process must send the message to an even or odd process rank.
- A solution is to create a communicator containing the even-ranked processes, another containing the odd-ranked processes, and initiate the collective communications inside these groups.

Communicators

Default communicator

- A communicator can only be created from another communicator. The first one will be created from the `MPI_COMM_WORLD`.
- After the `MPI_INIT()` call, a communicator is created for the duration of the program execution.
- Its identifier `MPI_COMM_WORLD` is an integer value defined in the header files.
- This communicator can only be destroyed via a call to `MPI_FINALIZE()`.
- By default, therefore, it sets the scope of collective and point-to-point communications to include all the processes of the application.

Communicators

Groups and communicators

- A communicator consists of :
 - A group, which is an ordered group of processes.
 - A communication context put in place by calling one of the communicator construction subroutines, which allows determination of the communication space.
- The communication contexts are managed by MPI (the programmer has no action on them : It is a hidden attribute).
- In the MPI library, the following subroutines exist for the purpose of building communicators : `MPI_COMM_CREATE()`, `MPI_COMM_DUP()`, `MPI_COMM_SPLIT()`
- The communicator constructors are collective calls.
- Communicators created by the programmer can be destroyed by using the `MPI_COMM_FREE()` subroutine.

Communicators

Partitioning of a communicator

In order to solve the problem example :

- Partition the communicator into odd-ranked and even-ranked processes.
- Broadcast a message inside the odd-ranked processes and another message inside the even-ranked processes.




Figure: Communicator creation/destruction

Communicators

Partitioning of a communicator with MPI_COMM_SPLIT()

The MPI_COMM_SPLIT() subroutine allows :

- Partitioning a given communicator into as many communicators as we want.
- Giving the same name to all these communicators : The process value will be the value of its communicator.
- Method :
 1. Define a colour value for each process, associated with its communicator number.
 2. Define a key value for ordering the processes in each communicator
 3. Create the partition where each communicator is called new_comm

```
1  MPI_COMM_SPLIT(comm,color,key,new_comm,code)
2
3  integer, intent(in) :: comm, color, key
4
5  integer, intent(out) :: new_comm, code
```

```
1  Comm.Split(self, int color=0, int key=0)
```

A process which assigns a color value equal to MPI_UNDEFINED will have the invalid communicator MPI_COMM_NULL for new_com.

Communicators

Example

Let's look at how to proceed in order to build the communicator which will subdivide the communication space into odd-ranked and even-ranked processes via the `MPI_COMM_SPLIT()` constructor.

Communicators

Partitioning of a communicator with MPI_COMM_SPLIT()

Example :

```
1  ...
2  world_rank = comm.Get_rank(); world_size = comm.Get_size()
3
4  m = 5; a = np.zeros(m)
5
6  if world_rank==2: a[:] = 2.
7  if world_rank==5: a[:] = 5.
8
9  key = world_rank
10 if world_rank==2 or world_rank==5 :
11     key=-1
12 color = world_rank%2
13
14 newcomm = comm.Split(color, key)
15
16 row_rank = newcomm.Get_rank(); row_size = newcomm.Get_size()
17
18 commname = "Comm-"+str(color)
19 newcomm.Set_name(commname)
20
21 print("WORLD RANK/SIZE: {RANK}/{SIZE} \t ROW RANK/SIZE: {ROW_RANK}/{ROW_SIZE} \t ↔
      Comm name: {commname}".format(RANK=world_rank, SIZE=world_size, ROW_RANK=↔
      row_rank, ROW_SIZE=row_size, commname=commname))
22
23 newcomm.Bcast(a, root=0);
24
25 newcomm.Free()
```

Communicators

Partitioning of a communicator with MPI_COMM_SPLIT()

Results :

```
mpirun -n 6 python3 split_communicator.py
```

WORLD RANK/SIZE: 5/6	ROW RANK/SIZE: 0/3	Comm name: Comm-1
WORLD RANK/SIZE: 0/6	ROW RANK/SIZE: 1/3	Comm name: Comm-0
WORLD RANK/SIZE: 1/6	ROW RANK/SIZE: 1/3	Comm name: Comm-1
WORLD RANK/SIZE: 2/6	ROW RANK/SIZE: 0/3	Comm name: Comm-0
WORLD RANK/SIZE: 3/6	ROW RANK/SIZE: 2/3	Comm name: Comm-1
WORLD RANK/SIZE: 4/6	ROW RANK/SIZE: 2/3	Comm name: Comm-0

Communicators

Communicator built from a group

- We can also build a communicator by defining a group of processes : Call to `MPI_COMM_GROUP()`, `MPI_GROUP_INCL()`, `MPI_COMM_CREATE()`, `MPI_GROUP_FREE()`
- This process is however far more cumbersome than using `MPI_COMM_SPLIT()` whenever possible.

Communicators

Topologies

- In most applications, especially in domain decomposition methods where we match the calculation domain to the process grid, it is helpful to be able to arrange the processes according to a regular topology.
- MPI allows defining virtual cartesian or graph topologies.
 - Cartesian topologies :
 - | Each process is defined in a grid.
 - | Each process has a neighbour in the grid.
 - | The grid can be periodic or not.
 - | The processes are identified by their coordinates in the grid.
 - Graph topologies :
 - | Can be used in more complex topologies.

Communicators

Cartesian topologies

- A Cartesian topology is defined from a given communicator named `comm_old`, calling the `MPI_CART_CREATE()` subroutine.
- We define :
 - An integer `ndims` representing the number of grid dimensions.
 - An integer array `dims` of dimension `ndims` showing the number of processes in each dimension.
 - An array of `ndims` logicals which shows the periodicity of each dimension.
 - A logical `reorder` which shows if the process numbering can be changed by MPI.

```
1 MPI_CART_CREATE(comm_old, ndims, dims, periods, reorder, comm_new, code↔  
    )  
2  
3 integer, intent(in) :: comm_old, ndims  
4 integer, dimension(ndims), intent(in) :: dims  
5 logical, dimension(ndims), intent(in) :: periods  
6 logical, intent(in) :: reorganization  
7  
8 integer, intent(out) :: comm_new, code
```

```
1 Intracomm.Create_cart(self, dims, periods=None, bool reorder=False)
```

Communicators

2D Example

Example on a grid having 2 domains along x and 2 along y, periodic in y.

```
1 from mpi4py import MPI
2
3 comm = MPI.COMM_WORLD
4 nb_procs = comm.Get_size()
5 rank = comm.Get_rank()
6
7 periods = tuple([False, False])
8 reorder = False
9 dims = [2,2]
10
11
12 cart2d = comm.Create_cart(
13     dims      = dims,
14     periods   = periods,
15     reorder   = reorder
16 )
```

- If `reorder = .false.` then the rank of the processes in the new communicator (`comm_2D`) is the same as in the old communicator (`MPI.COMM_WORLD`).
- If `reorder = .true.`, the MPI implementation chooses the order of the processes.




Figure: A 2D non-periodic Cartesian topology

Communicators

3D Example

Example on a 3D grid having 2 domains along x, 2 along y and 2 along z, non periodic.

```
1 from mpi4py import MPI
2
3 comm = MPI.COMM_WORLD
4 nb_procs = comm.Get_size()
5 rank = comm.Get_rank()
6
7 periods = tuple([False, False, False])
8 reorder = False
9 dims = [2,2,2]
10
11
12 cart3 = comm.Create_cart(
13     dims = dims,
14     periods = periods,
15     reorder = reorder
16 )
```






Figure: A 3D non-periodic Cartesian topology

Communicators

Process distribution

The `MPI_DIMS_CREATE()` subroutine returns the number of processes in each dimension of the grid according to the total number of processes.

```
1 MPI_DIMS_CREATE(nb_procs, ndims, dims, code)
2
3 integer, intent(in) :: nb_procs, ndims
4 integer, dimension(ndims), intent(inout) :: dims
5
6 integer, intent(out) :: code
```

Remark : If the values of `dims` in entry are all 0, then we leave to MPI the choice of the number of processes in each direction according to the total number of processes.

Communicators

Rank of a process

In a Cartesian topology, the `MPI_CART_RANK()` subroutine returns the rank of the associated process to the coordinates in the grid.

```
1 MPI_CART_RANK(comm, coords, rank, code)
2
3 integer, intent(in) :: comm
4 integer, dimension(ndims), intent(in) :: coords
5
6 integer, intent(out) :: rank, code
```

```
1 Cartcomm.Get_cart_rank(self, coords)
```

Communicators

Coordinates of a process

In a cartesian topology, the `MPI_CART_COORDS()` subroutine returns the coordinates of a process of a given rank in the grid.

```
1 MPI_CART_COORDS(comm, rank, ndims, coords, code)
2
3 integer, intent(in) :: comm, rank, ndims
4 integer, dimension(ndims), intent(out) :: coords
5
6 integer, intent(out) :: code
```

```
1 Cartcomm.Get_coords(self, int rank)
```

Communicators

Example : `MPI_CART_COORDS()`




Figure: A 2D non-periodic Cartesian topology

Communicators

Example : MPI_CART_COORDS()

```
1  from mpi4py import MPI
2
3  comm = MPI.COMM_WORLD
4  rank = comm.Get_rank()
5
6  periods = tuple([False, False])
7  reorder = False
8  dims = [2, 2]
9
10 cart2d = comm.Create_cart(
11     dims      = dims,
12     periods   = periods,
13     reorder   = reorder
14 )
15
16 coord2d = cart2d.Get_coords(rank)
17
18 print("I'm rank", rank, "my 2d coords are", coord2d)
```

`mpirun -n 4 python3 coordinate_2d_cart.py`

I'm rank 0 my 2d coords are [0, 0]

I'm rank 1 my 2d coords are [0, 1]

I'm rank 2 my 2d coords are [1, 0]

I'm rank 3 my 2d coords are [1, 1]

Communicators

Rank of neighbours

In a Cartesian topology, a process that calls the `MPI_CART_SHIFT()` subroutine can obtain the rank of a neighboring process in a given direction.

```
1 MPI_CART_SHIFT(comm, direction, step, rank_previous, rank_next, ↔  
    code)  
2  
3 integer, intent(in) :: comm, direction, step  
4 integer, intent(out) :: rank_previous, rank_next  
5  
6 integer, intent(out) :: code
```

```
1 Cartcomm.Shift(self, int direction, int disp)
```

- The direction parameter corresponds to the displacement axis (xyz).
- The step parameter corresponds to the displacement step.
- If a rank does not have a neighbor before (or after) in the requested direction, then the value of the previous (or following) rank will be `MPI_PROC_NULL`.

Communicators

Example : `MPI_CART_SHIFT()`




Figure: Call of the `MPI_CART_SHIFT()` subroutine

Communicators

Example : MPI_CART_SHIFT()

```
1 from mpi4py import MPI
2
3 comm = MPI.COMM_WORLD
4 rank = comm.Get_rank()
5
6 periods = tuple([False, False])
7 reorder = False
8 dims = [2, 2]
9
10 cart2d = comm.Create_cart(
11     dims = dims,
12     periods = periods,
13     reorder = reorder
14 )
15
16 left, right = cart2d.Shift(direction = 0, disp=1)
17 low, high = cart2d.Shift(direction = 1, disp=1)
18
19 print("I'm rank", rank, "my (left, right) neighbours are", (left, right),
20       "my (low, high) neighbours are", (low, high))
```

Communicators

Example : MPI_CART_SHIFT()

```
mpirun -n 4 python3 neighbours_2d_cart.py
```

```
I'm rank 0
```

```
my (left,right) neighbours are (-2, 2) my (low,high) neighbours are (-2, 1)
```

```
I'm rank 1
```

```
my (left,right) neighbours are (-2, 3) my (low,high) neighbours are (0, -2)
```

```
I'm rank 2
```

```
my (left,right) neighbours are (0, -2) my (low,high) neighbours are (-2, 3)
```

```
I'm rank 3
```

```
my (left,right) neighbours are (1, -2) my (low,high) neighbours are (2, -2)
```

Communicators

3D Example : coordinates and neighbours




Figure: 3D Coordinates and Neighbours

Communicators

3D Example : coordinates and neighbours

```
1 from mpi4py import MPI
2
3 comm = MPI.COMM_WORLD
4 rank = comm.Get_rank()
5
6 periods = tuple([False, False, False])
7 reorder = False
8 dims = [2,2,2]
9
10 cart3d = comm.Create_cart(
11     dims = dims,
12     periods = periods,
13     reorder = reorder
14 )
15
16 coord3d = cart3d.Get_coords(rank)
17 left, right = cart3d.Shift(direction = 0, disp=1)
18 low, high = cart3d.Shift(direction = 1, disp=1)
19 ahead, before = cart3d.Shift(direction = 2, disp=1)
20
21 print("I'm rank", rank, "my 3d coords are", coord3d, "my (left, right) neighbours ↔
      are", (left, right), "my (low, high) neighbours are", (low, high), "my (ahead↔
      , before) neighbours are", (ahead, before))
```

Communicators

3D Example : coordinates and neighbours

```
mpirun -n 8 --oversubscribe python3 create_3d_cart.py
```

```
I'm rank 0 my 3d coords are [0, 0, 0]  
my (left,right) neighbours are (-2, 4)  
my (low,high) neighbours are (-2, 2)  
my (ahead,before) neighbours are (-2, 1)
```

```
I'm rank 1 my 3d coords are [0, 0, 1]  
my (left,right) neighbours are (-2, 5)  
my (low,high) neighbours are (-2, 3)  
my (ahead,before) neighbours are (0, -2)
```

```
I'm rank 2 my 3d coords are [0, 1, 0]  
my (left,right) neighbours are (-2, 6)  
my (low,high) neighbours are (0, -2)  
my (ahead,before) neighbours are (-2, 3)
```

```
...
```